



Courage
Inspiration
Trust
Youth
Uniqueness

Welcome to City University of Hong Kong!

SDSC 3019 Lecture 9: P2P Architecture and Social Learning Network

Minghua Chen

<https://www.mhchen.com>

Announcement

- Project poster presentation in two weeks (Nov. 21) (20pts)
 - Submit a 2-minute video via Canvas by Nov 19, 8:59 am
 - Submit a poster via Canvas by Nov 19, 8:59 am
 - Arrangement to be announced on Canvas

- Project final report due on Dec. 2 (20 pts)
 - Up to 6 pages, IEEE Trans format

Outline

- Skype and BitTorrent: Peer-to-Peer Architectures
- Social Learning Networks: Optimizing interactions
- **Acknowledgement:** course materials are based on those shared by Prof. Brinton and Prof. Chiang at Purdue University

Skype and BitTorrent: Peer-to-Peer Architectures


Reading: Ch. 15

Outline


- The need for P2P applications
- P2P Skype and BitTorrent basics
- Overlay networks and tree graphs
- More on P2P Skype and BitTorrent
- Back-of-the-envelope download time
- Constructing multi-trees

Meeting challenges of scale


- Massive content distribution
 - Video and multimedia in particular
- Prevalent adoption of mobile wireless technologies
- How to scale up? Two approaches we will look at:
 - Help of peers: P2P (this lecture)
 - Using large data centers: Cloud (last lecture)
- Illustrates a key Internet principle:
The power of overlay
 - Under-specify protocols governing the operation of a network
 - Overlay network can be readily built on top of it for future applications



Server-based

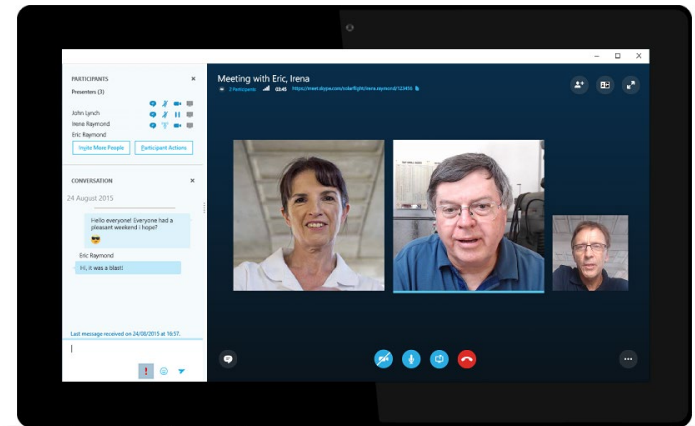


P2P-network



Peer-to-peer (P2P) applications

- Started becoming popular around 1999
 - Kazaa and Gnutella: Widely used P2P file- and music-sharing systems
 - Free riders: Consume much more than contribute
- Skype emerged in 2001
 - Bought by eBay in 2006 for \$2.6B, then Microsoft in 2011 for \$8B
 - 100M users worldwide as of 2020
- BitTorrent also started in 2001
 - 45M daily active users as of 2016
 - 20-40% of Internet traffic in mid 2010s



A screenshot of the Koinonein BitTorrent client interface. The top section shows a list of files being downloaded. The bottom section shows a detailed view of the peers connected to the selected file.

Name	% Done	Size	Uploaded	Upload Rate
ubuntu-17.10-desktop-amd64.iso	0%	1.39 GB	0 bytes	0 bytes/s
gimp-2.8.22-setup.exe	8%	85.4 MB	0 bytes	0 bytes/s
debian-9.2.1-amd64-netinst.iso	6%	290 MB	0 bytes	0 bytes/s

IP Address	Client	% Done	Uploaded	Upload Rate	Download Rate
178.70.240.167	µTorrent 3.5	100%	0 bytes	399 bytes/s	192 K
200.146.78.80	µTorrent 3.5	100%	0 bytes	111 bytes/s	358 K
89.87.69.138	µTorrent 3.5	100%	0 bytes	163 bytes/s	20 K
83.28.252.6	µTorrent 3.5	100%	0 bytes	101 bytes/s	210 K
80.99.27.18	µTorrent 3.5	100%	0 bytes	291 bytes/s	464 K
128.173.44.101	µTorrent Mac 1.8.7	100%	0 bytes	212 bytes/s	448 K
146.115.161.94	Transmission 2.92	100%	0 bytes	316 bytes/s	63.9 K


Skype vs. BitTorrent

- Skype *used* P2P for signaling
 - Locate each other, establish connections
- BitTorrent uses P2P for content sharing
 - Share chunks of files, leverage peer uplink capacities
- But both are free and scalable
- Both illustrate a positive network effect
 - Each additional node contributes to many other nodes
 - Can add more nodes without creating a bottleneck




Positive network effect

- Assumes nodes can effectively contribute
 - Requires some smart engineering designs
- **Metcalfe's law**
 - Benefit grows as the **square** of the number of nodes
 - Assumes all connections are equally important
 - Assumes each node is basically connected to all other nodes
- **P2P systems:** practical realization of this intuition
 - Achieve benefits by carefully choosing neighbors (even constant #)




BitTorrent basics

- Uses P2P for resource sharing
- Designed primarily for **multicasting**
 - Many users demanding the same file around the same time
 - Multicast vs. unicast
- Files divided into chunks, typically 256 kB
 - This way pieces can be shared
- Tracker: Centralized directory with peer-chunk database
 - Peer polls tracker for a set of 50 (or so) peers that has chunks it needs
 - Peer picks five of these 50 with which to exchange file chunks
- Peer refreshes its set of five peers at each timeslot
 - Partially based on which peers are helping the most
 - Partially based on randomization




Chunk vs. peering graph, client-server vs. P2P

- Each chunk traverses a tree of peers
 - Chunk is a rectangle, data transmission is dotted lines
 - A **tree** is an undirected graph with only one path from one node to any other node
 - Root node on top, leaf nodes on bottom
- Overall peering relationship is a general graph
 - Solid lines represent current peering relationships
 - Dotted lines represent potentials for the next time period
 - Evolves over time
- Data plane is distributed, P2P
 - In sharp contrast with the traditional **client-server** architecture





Chunk graph



Peer graph



Overlay network

- P2P is an **overlay network** $\tilde{G} = (\tilde{V}, \tilde{E})$
 - Underlay graph $G = (V, E)$
 - \tilde{V} is a subset of V
 - Links in \tilde{E} are some of the path in E
- Other overlays we've seen
 - The Internet is an overlay on top of physical networks
 - Online social networks are an overlay on top of the Internet
- The concept of an overlay is evolvable
 - Internet provides fundamental services
 - People can build overlay networks on top of existing ones
 - Example: IP multicast is rarely used; let IP take care of addressing, and instead use the P2P alternative with much less overhead




More on BitTorrent: Idea 1

- *Use smaller granularity than full files*
 - File chunks give more flexibility
- Transmission can be *spatially pipelined*
 - Each chunk can traverse a different multicast tree
 - Multi-tree transmission, multi-tree multicasting
- Compare this with packet switching
 - Divide messages into smaller granularity called packets
 - Multipath routing: Packets take possibly different paths
- “Peer”: Nodes are both senders and receivers of content




More on BitTorrent: Idea 2

- *Rarest chunk first strategy*
 - A peer should first download the chunks that very few peers have
- This optimizes content placement over different peers
- And equalizes the availability of chunks
 - Mitigates the problem where most of the peers have most of the chunks, but all must wait for the few rare chunks



More on BitTorrent: Idea 3

- *Peering construction method*
- First step: Tracker suggests 50-or-so potential peers to a new peer
- Recommendation based on:
 - The content they have
 - Performance-driven factors (e.g., distance to the new peer)
 - Peer churns: Which are still active
- Second step: Let the new peer pick, at each time, her actual “friends”
 - These are the peers to exchange chunks with
- Five peers are selected at each time slot:
 - Four are **tit-for-tat**: The top four peers in terms of the amount of content received
 - One is selected at random from the set of 50



Peer graph

Summary of BitTorrent Operation


- ① New peer A receives .torrent file from a web server
- ② Registers with tracker
- ③ Receives a list of neighbors
- ④ Selects five peers and establishes connections
- ⑤ Exchange bitmaps to indicate which chunks of content they have
- ⑥ With chunks selected, exchange chunks starting with the rarest ones
- ⑦ Every now and then, each peer updates its peer list

Client-server: Back-of-the-envelope calculation

- Consider N clients requesting a file of size F bits from a server
 - Server upload capacity u_s bps
 - Each of these clients has a download capacity of d_i bps
- How long will this take?
 - The server needs to send out NF bits, so it takes at least NF/u_s seconds
 - All the clients need to receive the file, including the slowest one with a download capacity of d_{min} , and that takes at least F/d_{min} seconds

So the total download time is:

$$T = \max \left\{ \frac{F}{d_{min}}, \frac{NF}{u_s} \right\}$$



- What happens as N becomes larger?
 - u_s cannot scale with N , especially as N becomes very large

P2P: Back-of-the-envelope calculation

- Suppose each peer i has an upload capacity u_i
 - In addition to a download capacity d_i as before
- Which tend to be larger: u_i or d_i ?
- Suppose we construct multicast trees that can fully utilize all u_i
 - For total number of bits to be shared, NF , the total upload bandwidth available to the whole network is $u_s + \sum_{i=1}^N u_i$
 - So the time this takes is $NF / (u_s + \sum_{i=1}^N u_i)$ seconds
- We also need to wait for ...
 - The server to send out each bit at least once: F / u_s seconds
 - The slowest peer to receive its bits: F / d_{min} seconds
- Time to distribute through the network:

$$T = \max \left\{ \frac{F}{u_s}, \frac{F}{d_{min}}, \frac{NF}{u_s + \sum_{i=1}^N u_i} \right\}$$

Back-of-the-envelope comparison

- Client-server:

$$T = \max \left\{ \frac{F}{d_{min}}, \frac{NF}{u_s} \right\}$$

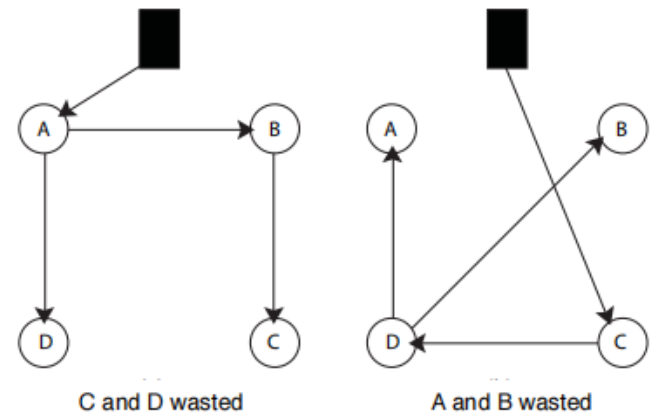
- P2P:

$$T = \max \left\{ \frac{F}{u_s}, \frac{F}{d_{min}}, \frac{NF}{u_s + \sum_{i=1}^N u_i} \right\}$$


- In both cases, last term has numerator increasing with N
- But for P2P, the denominator also increases with N
 - So, provided that the u_i are sufficiently large, T itself will no longer increase in N
 - The network performance scales itself with the network size

Constructing trees

- Back-of-the-envelope makes a huge assumption
 - All upload capacities can be fully utilized
 - But this can be hard to do, especially with only one tree
 - Consider each tree on the right: Some capacity not used in each case



- **Multi-tree** construction of peering relationships
 - Use more than one distribution tree
 - Helps, but the best way to construct all the trees is still not clear
 - Basic idea: Peers with more leftover upload capacities should be placed higher up in the next constructed trees
 - Still a very hard combinatorial optimization problem



Special case: No download bottleneck

- Consider the P2P case again, and assume the d_i are large enough
- We want to show that our back-of-the-envelope result can be obtained:

$$T = \max \left\{ \frac{F}{u_s}, \frac{NF}{u_s + \sum_{i=1}^N u_i} \right\}$$

- To do this, we can construct a multi-tree that collectively achieve the corresponding desired rates:


$$r_{max} = \min \left\{ u_s, \frac{u_s + \sum_{i=1}^N u_i}{N} \right\}$$

- We break this into two cases:

- ① $u_s \leq (u_s + \sum_{i=1}^N u_i)/N$
- ② $u_s > (u_s + \sum_{i=1}^N u_i)/N$

Case I: $u_s \leq (u_s + \sum_{i=1}^N u_i)/N$

- The maximum broadcast rate $r_{max} = u_s$ should be supported
 - Server upload capacity is too small
- Consider the following multi-tree of N trees, each two hops
 - Tree i has s take i as its child, and i take the other $N - 1$ peers as its children
 - Should deplete all upload capacity u_s



- Let each tree carry an upload capacity proportional to u_i :

$$r_i = \left(\frac{u_i}{\sum_{j=1}^N u_j} \right) u_s, \quad i = 1, \dots, N$$

Case I: $u_s \leq (u_s + \sum_{i=1}^N u_i)/N$

- This is possible because the upload speeds required are within capacity constraint. For the server,

$$\sum_{i=1}^N r_i = u_s$$

and for peer i ,

$$\begin{aligned}(N-1)r_i &= (N-1) \frac{u_i}{\sum_{j=1}^N u_j} u_s = \frac{u_i}{\sum_{j=1}^N u_j} (N u_s - u_s) \\ &\leq \frac{u_i}{\sum_{j=1}^N u_j} \left(u_s + \sum_{j=1}^N u_j - u_s \right) = u_i\end{aligned}$$

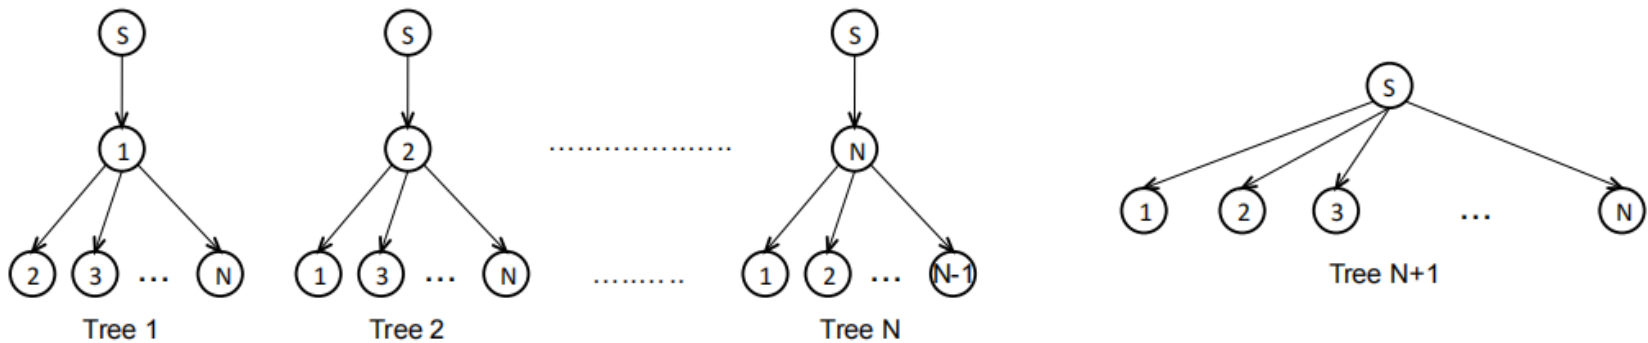
- The aggregate broadcast at which any peer i receives is then

$$r_{max} = r_i + \sum_{j \neq i} r_j = \sum_i r_i = u_s$$

- Hence, it takes F/u_s seconds to transfer the whole file

Case II: $u_s > (u_s + \sum_{i=1}^N u_i)/N$

- In this case, the maximum broadcast rate $r_{max} = (u_s + \sum_{i=1}^N u_i)/N$ should be supported
 - Server upload capacity is large enough for a different set of trees, including a server-client tree
- Consider the following multi-tree of $N + 1$ trees, N two hops and 1 one hop:



- Assume the trees carry the following rates:

$$r_i = \frac{u_i}{N-1} \quad i = 1, \dots, N \quad r_{N+1} = \frac{u_s - \frac{\sum_{i=1}^N u_i}{N-1}}{N}$$

meaning the last tree gets the leftover uplink capacity from the server

Case II: $u_s > (u_s + \sum_{i=1}^N u_i)/N$

- These upload speeds required are within constraint. For peer i ,

$$(N - 1)r_i = u_i$$

and for the sever,

$$\begin{aligned}\sum_{i=1}^N r_i + N \cdot r_{N+1} &= \sum_{i=1}^N \frac{u_i}{N-1} + N \cdot \frac{u_s - \frac{\sum_{i=1}^N u_i}{N-1}}{N} \\ &= \sum_{i=1}^N \frac{u_i}{N-1} + u_s - \frac{\sum_{i=1}^N u_i}{N-1} = u_s\end{aligned}$$

- The aggregate broadcast at which any peer i receives is then

$$\begin{aligned}r_{max} &= r_i + r_{N+1} + \sum_{j \neq i} r_j = \frac{u_i}{N-1} + \frac{u_s - \frac{\sum_{i=1}^N u_i}{N-1}}{N} + \sum_{j \neq i} \frac{u_j}{N-1} \\ &= \frac{N \sum_i u_i}{N(N-1)} + \frac{u_s}{N} - \frac{\sum_{i=1}^N u_i}{N(N-1)} = \frac{u_s + \sum_i u_i}{N}\end{aligned}$$

Case II: $u_s > (u_s + \sum_{i=1}^N u_i)/N$

- Hence, it takes $NF/(u_s + \sum_{i=1}^N u_i)$ seconds to transfer the whole file
- Combining the two cases above produces our desired results






Outline

- Online learning
- Social Learning Networks (SLN)
- Learner benefit in an SLN
- SLN utility maximization
- Numerical example
- Summary

Distance learning

- Mid-1800s: Degree programs operating by postal mail
- Early-to-mid-1900s: Course lecture broadcasts over radio and television
- 1990s: Online programs, degrees, and universities
 - Rise of the web
 - By 2003: 80% of colleges had at least one class using eLearning technology
 - By 2014: More than 95% of public colleges offered some form of online program

- **Learning modes**
offered by different
technologies:

	Auditory 	Visual 	Textual 	Social 	Synchronous 
In Class	✓	✓	✓	✓	✓
Mail		✓	✓		
Radio	✓				
Television	✓	✓			
Internet	✓	✓	✓	✓	

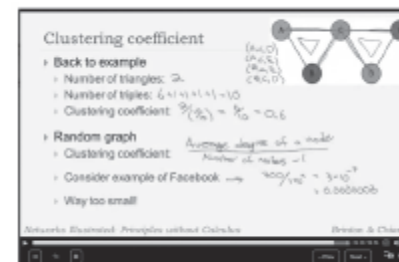
Massive Open Online Courses (MOOCs)

- ❑ 2002: MIT creates an online resource of free course materials
 - Anyone can access
 - Why would MIT do this?
- ❑ Became a pathfinder for **MOOC**
 - Transform classroom-sized online courses into MOOCs
 - Coursera, edX, Udacity, Udemy, ...
 - Partnership with top universities
- ❑ MOOCs are ...
 - Massive: Up to hundreds of thousands per session
 - Open: Typically free or very cheap
 - Online
 - Course: Typically full course delivery

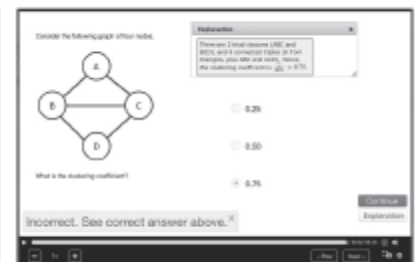
coursera



U UDACITY





Lecture video



In-video quiz


Social Learning Networks (SLN)

- Networks of ...
 - Students
 - Instructors
 - Modules of learning
- One of the scalable aspect of MOOCs
 - Address diverse sets of learning needs
 - Interesting spin on how to achieve individualization
- Three aspects
 - *Social* : Hinges on interaction between peers
 - *Network*: The topology (student-to-student)
 - *Learning*: The functionality
- How to study and optimize SLN?
 - Massive amounts of data collected




Data Sources in eLearning

- ❑ At least four groups of data:
 - ❑ Clickstream measurements
 - Learners watching video lectures
 - ❑ Social interactions
 - Learners on discussion forums
 - ❑ Question responses
 - How knowledge transfer is normally assessed
 - ❑ Content (and user) words
 - Topic-specific learning
 - ❑ How to use these to analyze and optimize SLN?
 - Methodologies at the intersection of data science, networking, and learning sciences


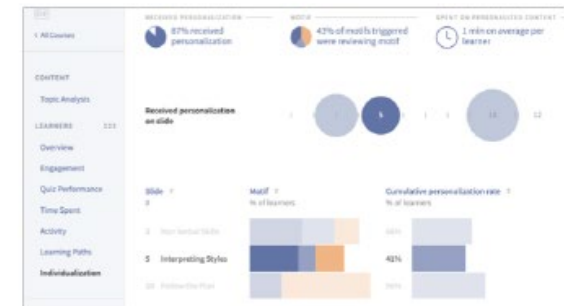
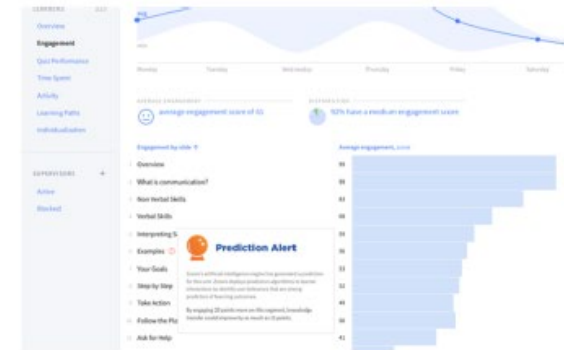


3. Find x .




Zoomi: Artificial Intelligence for Learning

- Predictive Learning Analytics (PLA)
 - Fine granular data analysis
 - Behavior-based early detection outcome prediction
 - Behavioral pattern identification
- Network Efficiency Optimization (NEO)
 - Predicting link formation over time
 - Optimizing interaction structures (our focus in this lecture)
 - Comparing SLN and course learning outcomes
- Deep Learning Personalization (DLP)
 - Autonomous creation of remediation content
 - Reinforcement learning-based personalization algorithms




Graph of an SLN

- Nodes: Learners sharing information
 - Fine granular data analysis
 - Index users by $u \in U$, the set of all users
- Links: Interaction structure between users
 - Define $\mathbf{W} = [w_{u,v}]$ as the weighted adjacency matrix
 - $w_{u,v}$ represents spread of information from u to v
- How to define/interpret $w_{u,v}$?
 - Several possibilities
 - One useful designation: Consider $w_{u,v} \in [0,1]$ as the probability that u responds when v posts
 - **\mathbf{W} is asymmetric**, i.e., $w_{u,v} \neq w_{v,u}$ in general





Modeling SLN Discussions

- Topics: Focal points of discussions
 - These can be latent dimensions
 - Index by $k \in K$, the set of discussion topics
- Seeking: Tendency to ask questions
 - Topic specific Let $S = [s_{u,k}]$ be the matrix of seeking tendencies, with $s_{u,k} \geq 0$ being u 's seeking tendency on topic k
- Disseminating: Tendency to provide answers/facts about the material
 - Similarly topic specific
 - Let $D = [d_{u,k}]$ be the matrix of disseminating tendencies, with $d_{u,k} \geq 0$ being u 's disseminating tendency on topic k



Modeling User Benefits

- Learning benefit: Incoming weights on topics with high seeking tendency
 - We quantify this as: $s_{u,k} \cdot f(\sum_v w_{u,v} d_{v,k})$
 - $w_{u,v} d_{v,k}$ is the expected amount of response from v to u on topic k
 - f is a concave function: Diminishing marginal returns
 - Weighted by seeking tendency $s_{u,k}$
- Teaching benefit: Outgoing weights on topics with high disseminating tendency
 - Similarly, this can be quantified as:
 $d_{u,k} \cdot f(\sum_v w_{u,v} s_{v,k})$
 - $w_{u,v} s_{v,k}$ is the amount by which u provides information to v sought on topic k
 - f is a concave function
 - Weighted by $d_{u,k}$



Quantifying Utility

- User-topic benefits: $b_{u,k} \geq 0$ is the total utility obtained by u with respect to k , quantified as

$$b_{u,k} = s_{u,k} \log(1 + \sum_v w_{v,u} d_{v,k}) + \alpha_u \cdot d_{u,k} \log(1 + \sum_v w_{u,v} s_{v,k}),$$

In matrix form, $\mathbf{B} = [b_{u,k}]$

- α_u quantifies the benefit of teaching relative to learning for user u
- Choose $f(x) = \log(1 + x)$ as a standard concave function

- Seeking to Incoming Disseminating Ratio (SIDR):

$$\phi_{u,k} = \frac{s_{u,k}}{\sum_v w_{v,u} d_{v,k}}, \quad \Phi = [\phi_{u,k}]$$

- Lower makes user's learning benefit in $b_{u,k}$ higher
- More concretely, smaller SIDR $\phi_{u,k}$ implies that u 's seeking tendency on topic k has higher satisfaction from the incoming disseminating tendencies of their neighbors

Quantifying Utility

- Disseminating to Incoming Seeking Ratio (DISR):

$$\psi_{u,k} = \frac{d_{u,k}}{\sum_v w_{u,v} s_{v,k}}, \quad \Psi = [\psi_{u,k}]$$

- Lower makes user's teaching benefit in $b_{u,k}$ higher
- A smaller DISR $\psi_{u,k}$ implies that u 's disseminating tendency on k is being used to satisfy more of their neighbor's seeking tendency

- Local utility (total benefit obtained by u across k):

$$l_u = \sum_k b_{u,k}$$

- Global utility (average local utility across users):

$$g_{\alpha_u} = \frac{1}{|\mathcal{U}|} \sum_{u,k} b_{u,k}$$

- Can we maximize g without sacrificing l ?, by tuning

Optimizing the SLN

- Formulate as a Network Utility Maximization (NUM) problem:

$$\begin{array}{ll}
 \underset{\mathbf{W}}{\text{maximize}} & g(\mathbf{W}) \\
 \text{subject to} & \Phi(\mathbf{W}) \leq C_s \hat{\Phi} \\
 & \Psi(\mathbf{W}) \geq C_d \hat{\Psi} \\
 & \mathbf{0} \leq \mathbf{W} \leq \mathbf{1}, \text{diag}(\mathbf{W}) = \mathbf{0} \\
 \text{variables} & \mathbf{W}
 \end{array}$$

First two constraints:

- Preserving incoming information: $\Phi(W) \leq C_s \hat{\Phi}$
 - $\hat{\Phi} = [\hat{\phi}_{u,k}]$ is the matrix of observed SIDR from the observed $\hat{\mathbf{W}}$
 - $C_s > 0$ is a parameter: If $C_s < 1$, we require a better SIDR than before, while if $C_s > 1$, we provide some slack to the feasible set
- Balancing load: $\Psi(W) \geq C_d \hat{\Psi}$
 - $\hat{\Psi} = [\hat{\psi}_{u,k}]$ is the matrix of observed SIDR
 - $C_d > 0$: If $C_d > 1$, we require users to participate more, while if $C_d < 1$, we allow any particular DISR to drop if needed

Quantifying the Impact of Optimization

- How to quantify the effect of optimization?
 - Several dimensions
 - Can boil down to a comparison of network **efficiency** and **fairness**
- **Efficiency**: Fraction of optimal global utility achieved in the observed network, i.e.,

$$\eta_{\alpha_u, C_s, C_d} = g_{\alpha_u}(\hat{W}) / g_{\alpha_u, C_s, C_d}^*$$

- **Fairness**: Ratio of fairness after vs. before, i.e.,

$$\epsilon_{\alpha_u, C_s, C_d} = F(\hat{\mathbf{l}}_{\alpha_u}(\hat{W})) / F(\mathbf{l}_{\alpha_u, C_s, C_d}^*)$$

- $\mathbf{l} = (l_1, l_2, \dots)$ is the vector of local utilities, $\hat{\mathbf{l}}$ is before and \mathbf{l}^* is after optimization
- $F()$ is a function quantifying fairness of a distribution, e.g., α -fairness (from data pricing) or Jain's Index

Inference and Optimization Algorithms

□ Quantifying the observed network \widehat{W}

- If v posts in a thread, what is the probability u will post after that?

$$\hat{w}_{u,v}(\sigma) = \frac{n_{u,v} + \sigma(\sum_j n_{u,j} / \sum_j N_j) N_{max}}{N_v + \sigma N_{max}}$$

- $n_{u,v}$ is the number of times that u posts after v , N_v is the total number of times v posted
- σ is a Bayesian adjustment parameter because different users post different amounts (remember the product rating lecture)

□ Inferring seeking and disseminating tendencies

- Obtain post-topic distributions through e.g., Latent Dirichlet Allocation, a common topic modeling procedure
- Label posts as question or answer through e.g., rule-based information retrieval methods

□ Solving for the optimal network W^* ?

- Convex optimization problem (convex objective, linear constraints)
- But it is a very big problem: W can have millions of entries

Numerical Example: Dataset

- Give $\hat{\mathbf{W}}$, \mathbf{S} , and \mathbf{D} matrices for a toy SLN with $|\mathbf{U}|=10$ users and $|\mathbf{K}|=10$ topics, we would have to compute the $s_{u,k}$ and $d_{u,k}$ based on e.g., the discussion forum structure.
- Here are some of the matrix values:

$$\hat{\mathbf{W}} = \begin{bmatrix} 0.0 & 0.0 & 0.0 & 0.0 & \dots \\ 0.1 & 0.0 & 0.0 & 0.0 & \dots \\ 0.1 & 1.0 & 0.0 & 0.0 & \dots \\ 0.1 & 1.0 & 0.166 & 0.0 & \dots \\ \vdots & \vdots & \vdots & \vdots & \vdots \end{bmatrix} \quad \mathbf{S} = \begin{bmatrix} 0.0 & 0.0 & 0.0 & 0.0 & \dots \\ 0.0 & 0.0 & 0.0 & 0.0 & \dots \\ 0.854 & 1.192 & 0.390 & 0.116 & \dots \\ 1.542 & 0.042 & 0.860 & 0.042 & \dots \\ \vdots & \vdots & \vdots & \vdots & \vdots \end{bmatrix}$$
$$\mathbf{D} = \begin{bmatrix} 5.538 & 2.821 & 0.691 & 2.429 & \dots \\ 1.754 & 0.529 & 0.046 & 0.046 & \dots \\ 0.798 & 0.107 & 0.794 & 0.107 & \dots \\ 0.322 & 0.041 & 0.322 & 0.041 & \dots \\ \vdots & \vdots & \vdots & \vdots & \vdots \end{bmatrix}$$

Numerical Example: DISR and SIDR

- First, we calculate the DISR ψ and SIDR ϕ values. For example, $\phi_{4,1}$

$$\begin{aligned}\hat{\psi}_{4,1} &= \frac{d_{4,1}}{\sum_v w_{4,v} s_{v,1}} = \frac{0.322}{0.1 \times 0.0 + 1.0 \times 0.0 + 0.166 \times 0.854 + \dots} \\ &= 2.26\end{aligned}$$

- So user $u = 4$ has about twice as much dissemination on topic $k = 1$ than what is being sought by her neighbors. As another example, $\phi_{3,4}$ is:

$$\begin{aligned}\hat{\phi}_{3,4} &= \frac{s_{3,4}}{\sum_v w_{v,3} d_{v,4}} = \frac{0.116}{0.0 \times 2.429 + 0.0 \times 0.046 + 0.0 \times 0.107 + \dots} \\ &= 0.0097\end{aligned}$$

- So user $u = 3$ has substantially more incoming dissemination on topic $k = 4$ than what she is asking for.

Numerical Example: DISR and SIDR

- Continuing for all (u; k) pairs, the observed DISR and SIDR matrices are ($\phi_{u,k} = \infty$ means $d_{u,k} > 0$ while $\sum_v w_{u,v} s_{v,k} = 0$):

$$\hat{\Psi} = \begin{bmatrix} 146.9 & 73.5 & 7.8 & 49.6 & \dots \\ \infty & \infty & \infty & \infty & \dots \\ 13.5 & 0.9 & 13.6 & 1.2 & \dots \\ 2.2 & 0.2 & 4.9 & 2.1 & \dots \\ \vdots & \vdots & \vdots & \vdots & \vdots \end{bmatrix} \quad \hat{\Phi} = \begin{bmatrix} 0.0 & 0.0 & 0.0 & 0.0 & \dots \\ 0.0 & 0.0 & 0.0 & 0.0 & \dots \\ 0.04 & 0.10 & 0.010 & 0.009 & \dots \\ 0.06 & 0.003 & 0.020 & 0.003 & \dots \\ \vdots & \vdots & \vdots & \vdots & \vdots \end{bmatrix}$$

- We can see, for example:
 - users 1 and 2 are not asking questions ($\hat{\phi}_{1,k}, \hat{\phi}_{2,k} = 0$) but have a substantial amount of dissemination to offer for those they are not connected to ($\hat{\psi}_{1,k}, \hat{\psi}_{2,k} \gg 0$ while $\hat{w}_{1,v}, \hat{w}_{2,v} = 0$ for most v)
 - most users have substantial incoming dissemination ($\hat{\phi}_{u,k} \ll 1$) based on the amount of questions they are asking

Numerical Example: Benefits and Utilities

- We can also calculate the observed benefits $\hat{b}_{u,k}$, local utilities \hat{l}_u , and global utility \hat{g} (the u for each user are provided with the CSVs as well). For user 3 on topic 1, the benefit is

$$\begin{aligned} b_{3,1} &= s_{3,1} \log(1 + \sum_v w_{v,3} d_{v,1}) + \alpha_u \cdot d_{3,1} \log(1 + \sum_v w_{3,v} s_{v,1}) \\ &= 0.85431 \cdot \log(1 + (0 \times 5.53 + 0 \times 1.75 + 0 \times 0.798 + \dots)) \\ &\quad + 0.03421 \cdot 0.7983 \cdot \log(1 + (0.1 \times 0 + 1 \times 0 + \dots)) = 1.1380 \end{aligned}$$

with $u = 0.03421$. The local utility for this user is

$$\hat{l}_3 = b_{3,1} + b_{3,2} + b_{3,3} + \dots = 10.3716$$

Repeating this for every user, the local utilities are

$$\hat{\mathbf{l}} = (0.09, 0, 4.27, 10.37, 20.78, 0.06, 0.009, 17.93, 0.010, 0.077)$$

So users 4, 5, and 8 have the most utility. The global utility is

$$\hat{g} = (1/10) \sum \hat{l}_u = 5.3596$$

Numerical Example: Benefits and Utilities

- An optimal solution \mathbf{W}^* for this problem is provided in the CSVs also. A few of the entries are as follows:

$$\mathbf{W}^* = \begin{bmatrix} 0.0 & 0.0 & 0.0 & 0.0 & \dots \\ 0.053 & 0.0 & 0.0 & 0.0 & \dots \\ 0.079 & 0.956 & 0.0 & 0.00061 & \dots \\ 0.068 & 1.0 & 0.166 & 0.0 & \dots \\ \vdots & \vdots & \vdots & \vdots & \vdots \end{bmatrix}$$

- Compared with $\widehat{\mathbf{W}}$, for example, we see that $w_{2,1}^* \approx \widehat{w}_{2,1}$, i.e., user 2 should respond to user 1 with about half the frequency. Also, while, $\widehat{w}_{3,4} = 0$, $w_{3,4}^* > 0$ i.e., user 3 should respond to user 4's questions (albeit infrequently).
- With this, we can find the optimal local and global utilities to be:

$$\mathbf{I}^* = (0.12, 0.0003, 4.28, 10.95, 24.16, 1.79, 0.009, 19.66, 0.013, 0.095)$$

and $g^* = 6.11$.

Numerical Example: Comparing Observed and Optimal

- Comparing \hat{g} to g^* , the efficiency is

$$\eta = 5.36/6.11 = 0.88,$$

i.e., there is about 12% improvement under optimization. To compare $\hat{\mathbf{l}}$ to \mathbf{l}^* , we need to choose an efficiency measure. For Jain's Index, the fairness of a vector $\mathbf{x} = (x_1, x_2, \dots, x_n)$ is

$$F(\mathbf{x}) = \frac{(\sum_{i=1}^n x_i)^2}{n \cdot \sum_{i=1}^n x_i^2}$$

- With this, $F(\hat{\mathbf{l}}) = 0.327$ and $F(\mathbf{l}^*) = 0.336$, which gives a fairness ratio of

$$\epsilon = 0.327/0.336 = 0.97$$

- So, the fairness actually improves (albeit slightly) under optimization.



End-of-Lecture Quiz

and

Weekly Anonymous Survey

closed book, closed notes, and no discussion

Thank you!




Minghua Chen (minghua.chen@cityu.edu.hk)

<http://personal.cityu.edu.hk/~mchen88/>

P2P Skype basics

- Two major technologies
 - Voice over IP (VoIP)
 - P2P: Our focus here
- Phone calls are intrinsically P2P, but Skype uses P2P to:
 - Discover peers
 - Traverse **firewalls**: Software and hardware blocking incoming data connections
- **Super Nodes (SNs)**: Machines with public IP addresses
 - Contain replicas of Skype's central directory (including IP)
 - Distributed throughout the network, forming a full mesh
 - Act as a network of publicly visible relays
- What happens when caller and callee are behind firewalls?
 - Caller initiates connection with an SN, and callee with another SN
 - Two-way communication established via the SNs
 - Mutually agree on a single SN to shorten the path



More on P2P Skype

- Super nodes (SN)
 - Needs public IP address to traverse NATs and firewalls
 - Should have abundant resources: CPU, memory, ...
- Two tiers of Skype overlay
 - Mesh of SNs
 - Shallow tree of ordinary hosts rooted at each SN
- During login, a host ...
 - Advertises its presence to other hosts
 - Discovers whether it is behind a firewall, which SN to connect to, and which public IP address to use

